

Aarush Agarwal

aarusha@andrew.cmu.edu | linkedin.com/in/agarwal-aarush | github.com/AgarwalAarush | aarushagarwal.dev

EDUCATION

Carnegie Mellon University

Bachelor of Science in Artificial Intelligence & Computer Science

Expected December 2027

3.78 QPA, Dean's List High Honors

Honors: USACO Gold, 6× Hackathon Winner (Claude, Harvard, Berkeley, 3× CMU)

EXPERIENCE

Shopify | Machine Learning Engineer Intern, Search Relevance

Jun 2026 - Present

- Designed a merchant-aware auxiliary training objective that **improved a core merchant-search relevance metric by 10%**, prioritizing first-party catalogs over third-party resellers for merchant-intent queries. Enabled the team to incorporate high-quality clicks as secondary supervision alongside relevance labels.
- Fine-tuned, distilled, and deployed Liquid AI LFM2.5 re-rankers with Triton Inference Server and PyTorch Inductor, **cutting production p99 latency by 5%** (20 → 19 ms) via CUDA graphs and tokenization optimizations.
- Identified data/supervision mismatches through win/loss analysis of production re-rankers vs. LLM reference judgments, establishing the need for new training data.
- Developing a unified multilingual query-rewriting model, improving an internal search-relevance metric by 0.5%.

Shopify | Machine Learning Engineer Intern, Buyer Fraud

May 2025 - Aug 2025

- Co-filed a patent** on a distributed multi-agent system that decomposes tasks via a Neo4j dependency graph and executes subtasks across specialized agents in parallel.
- Improved buyer-fraud detection accuracy by 3%** and reduced training iteration time by 70% through dimensionality reduction, importance-based feature pruning, and BigQuery/Dataflow + Vertex AI pipeline rebuilds.

CMU Language Technologies Institute | ML Researcher

Jan 2026 - Present

- Researching adaptive multimodal MoE architectures with Prof. Chenyan Xiong, targeting capacity expansion for OOD vision-language workloads while mitigating text-reasoning degradation during continual pretraining and fine-tuning.
- Building an autonomous MoE experimentation framework that uses specialist agents to propose bounded architecture, routing, curriculum, and adapter changes; evaluating variants against baselines using multimodal loss proxies, router-health checks, and lm-eval reasoning benchmarks.

CMU Cosmology Lab & CERN | CUDA Researcher

Aug 2024 - Oct 2025

- First-authored [FastGraph](#)**, a GPU-resident exact k-nearest-neighbor engine for low-to-moderate-dimensional graph construction, delivering up to **41× speedups over FAISS-GPU and GGNN**. Developed PCA-subspace spatial pruning that reduces the search space while evaluating distances in the original coordinates, preserving exact neighbors.
- Engineered compile-time-specialized C++/CUDA kernels with 2–5D spatial binning, hypercube neighborhood traversal, register-cached distance evaluation, custom backward gradients, and PyTorch/GravNet integration for differentiable graph-learning workflows.

PROJECTS

AutoReflex | Hardware & ML

Mar 2026

- 3rd Place at HackberryPi** for AutoReflex, a hardware-in-the-loop control system combining 100 FPS global shutter vision on a Jetson Orin Nano, UDP coordinate streaming, and a 1 kHz Raspberry Pi PID loop to drive a servo, solenoid, and TENS actuation in approximately 30 ms end to end.

Haven | Backend

Oct 2025

- Won 1st Place in the YC, LiveKit, and Fetch.AI tracks and Best Startup at CalHacks** for a patient-care platform; architected clinical-agent infrastructure, an rPPG heart-rate monitoring pipeline (OpenCV, FastICA), and real-time voice intake (LiveKit, Groq).

SKILLS

Languages Python, C++, CUDA, SQL, C, Swift

Software PyTorch, TensorFlow, JAX, OpenCV, GCP, PostgreSQL, Neo4j, Triton Inference Server, TensorRT

ML Systems Model Distillation, Distributed Training, Mixture-of-Experts, Computer Vision, Search Ranking